

```
In [1]: #import necessary Libraraies
import numpy as np #linear algebra
import pandas as pd #data processing
import matplotlib.pyplot as plt #plots
%matplotlib inline
import seaborn as sns
```

```
In [3]: #read in the data
tr = pd.read_csv('Train (1).csv')
te = pd.read_csv('Test (1).csv')
ss = pd.read_csv('SampleSubmission.csv')

#make a copy of the data
train = tr.copy()
test = te.copy()
```

```
In [5]: !pip install catboost
```

Collecting catboost

Downloading https://files.pythonhosted.org/packages/90/86/c3dcb600b4f9e7584ed90ea9d30a717fb5c0111574675f442c3e7bc19535/catboost-0.24.1-cp36-none-manylinux1_x86_64.whl (https://files.pythonhosted.org/packages/90/86/c3dcb600b4f9e7584ed90ea9d30a717fb5c0111574675f442c3e7bc19535/catboost-0.24.1-cp36-none-manylinux1_x86_64.whl) (66.1MB)

 66.1MB 84kB/s

Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from catboost) (1.4.1)

Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from catboost) (4.4.1)

Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.1.2)

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from catboost) (1.15.0)

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.18.5)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from catboost) (3.2.2)

Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages (from catboost) (0.10.1)

```
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from plotly->catboost) (1.3.3)
```

Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2.8.1)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2018.9)

Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (0.10.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (1.2.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (2.4.7)

```
Installing collected packages: catboost
```

Successfully installed catboost-0.24.1

```
In [6]: #display all rows and columns
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)
```

EDA

```
In [ ]: train.head() #veiw the dataframe
```

```
Out[6]:
```

	Applicant_ID	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form_f
0	Apcnt_1000000	3436.0	0.28505	1.6560	0.0	0.000	0.0	10689
1	Apcnt_1000004	3456.0	0.67400	0.2342	0.0	0.000	0.0	898
2	Apcnt_1000008	3276.0	0.53845	3.1510	0.0	6.282	NaN	956
3	Apcnt_1000012	3372.0	0.17005	0.5050	0.0	0.000	192166.0	3044
4	Apcnt_1000016	3370.0	0.77270	1.1010	0.0	0.000	1556.0	214

```
In [ ]: train.isnull().sum() #check for missing values
```

```
Out[7]: Applicant_ID      0
form_field1      2529
form_field2      3844
form_field3      355
form_field4      355
form_field5      355
form_field6     13360
form_field7      5163
form_field8     13360
form_field9      8008
form_field10     355
form_field11     31421
form_field12     9895
form_field13     5889
form_field14      0
form_field15     22475
form_field16     13036
form_field17     11151
form_field18     10402
```

```
In [ ]: train.describe() #data description of each column
```

```
Out[8]:
```

	form_field1	form_field2	form_field3	form_field4	form_field5	form_field6	form
count	53471.000000	52156.000000	55645.000000	55645.000000	55645.000000	4.264000e+04	5.0837
mean	3491.795665	0.550737	1.052225	0.851979	1.956317	6.244479e+05	6.8652
std	188.462426	0.820979	2.147768	3.157692	10.512396	1.433422e+06	1.9127
min	2990.000000	0.000000	0.000000	0.000000	0.000000	0.000000e+00	0.0000
25%	3358.000000	0.070788	0.000000	0.000000	0.000000	1.400400e+04	6.8697
50%	3484.000000	0.267575	0.062000	0.000000	0.000000	1.155330e+05	2.7043
75%	3620.000000	0.719512	1.282000	0.000000	0.000000	5.259280e+05	6.9938
max	3900.000000	18.015050	57.371600	91.672200	407.748600	5.313546e+07	2.1587

```
In [7]: #encode categorical columns
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
train['form_field47'] = le.fit_transform(train.form_field47)
test['form_field47'] = le.fit_transform(test.form_field47)
train['default_status'] = le.fit_transform(train.default_status)
```

```
In [8]: features_with_integer_values = ['form_field16', 'form_field17', 'form_field18',
                                         'form_field19', 'form_field20', 'form_field34', 'fo
                                         'form_field36', 'form_field37', 'form_field38', 'fo

features_recorded_in_days = ['form_field26', 'form_field27', 'form_field28', 'form_

features_recorded_in_months = ['form_field32']

categorical_features = ['form_field47']

#All other features are of float type
```

```
In [9]: mean_fill = features_with_integer_values + features_recorded_in_days + features_
#mean_fill
```

```
In [10]: #fill the missing values for these columns with their mean
for i in mean_fill:
    train[i].fillna(train[i].mean(), inplace = True)
    test[i].fillna(test[i].mean(), inplace = True)
```

```
In [11]: cont_feat = ['form_field1', 'form_field2', 'form_field3', 'form_field4', 'form_field
                    'form_field9', 'form_field10', 'form_field11', 'form_field12', 'form_fie
                    'form_field22', 'form_field23', 'form_field24', 'form_field25', 'form_f
                    'form_field42', 'form_field43', 'form_field44', 'form_field48', 'form_fi
```

```
In [12]: #fill the missing values for these columns with their -1
for i in cont_feat:
    train[i].fillna(-1, inplace = True)
    test[i].fillna(-1, inplace = True)
```

```
In [13]: from sklearn.preprocessing import StandardScaler
```

```
In [14]: #drop unwanted column
train.drop('Applicant_ID', axis=1, inplace= True)
test.drop('Applicant_ID', axis=1, inplace= True)
```

```
In [15]: #instantiate Standard scaler
sc = StandardScaler()

#scale the data... scaling is done to center the data distribution around zero and
scaled_train = pd.DataFrame(sc.fit_transform(train), columns = train.columns)
scaled_test = pd.DataFrame(sc.fit_transform(test), columns = test.columns)
```

```
In [16]: #divide the data to features and label
X = train.drop('default_status', axis = 1)
y = train['default_status']
```

MODELLING

```
In [17]: #import models
from sklearn.model_selection import StratifiedKFold, RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier # algorithm for training an ensemble
from xgboost import XGBClassifier # extreme boosting algorithm
from lightgbm import LGBMClassifier # extreme boosting algorithm
from catboost import CatBoostClassifier
from sklearn.metrics import roc_curve, auc, roc_auc_score
```

```
In [22]: # instantiate the models
rfc = RandomForestClassifier(n_estimators=400, random_state=42, n_jobs= -1)
xgb = XGBClassifier(n_estimators=300, random_state=42, scale_pos_weight=3.08, min_child_weight=1)
lgb = LGBMClassifier(reg_lambda = 2, random_state = 42, n_estimators = 4000)
cb = CatBoostClassifier(n_estimators=4000, eval_metric='AUC', learning_rate=0.1, max_depth=8,
                        bootstrap_type='Bayesian', use_best_model=True, od_wait=50)
```

```
In [19]: #stratifiedKFold parameters
n=10
skf = StratifiedKFold(n_splits = n, shuffle = True, random_state=42)
```

CATBOOST MODEL

```
In [23]: scores_cb = []
pred_test1 = 0
for train_index, test_index in skf.split(X,y):
    print('train index: ', train_index, '\n')
    print('test index: ', test_index, '\n')
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    cb.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_test, y_test)], verbose=0)
    pred = cb.predict_proba(X_test)[:,1]
    print('score: ', roc_auc_score(y_test, pred))
    scores_cb.append(roc_auc_score(y_test, pred))
    test_pred1 = cb.predict_proba(test)[:,1]
    pred_test1 += test_pred1
```

```
train index: [ 0 1 2 ... 55996 55997 55999]
```

```
test index: [ 11 56 58 ... 55983 55984 55998]
```

```
0:      test: 0.7835423 test1: 0.7990652      best: 0.7990652 (0)      tota
1: 33.6ms      remaining: 2m 14s
100:    test: 0.8390957 test1: 0.8502515      best: 0.8502515 (100)    tota
1: 3.41s      remaining: 2m 11s
200:    test: 0.8461846 test1: 0.8529336      best: 0.8530022 (187)    tota
1: 6.76s      remaining: 2m 7s
300:    test: 0.8512596 test1: 0.8532851      best: 0.8533724 (241)    tota
1: 10.1s      remaining: 2m 4s
400:    test: 0.8559323 test1: 0.8538175      best: 0.8538424 (398)    tota
1: 13.5s      remaining: 2m 1s
500:    test: 0.8598518 test1: 0.8537553      best: 0.8538646 (402)    tota
1: 16.9s      remaining: 1m 57s
Stopped by overfitting detector (120 iterations wait)
```

```
bestTest = 0.8538646099
```

```
In [24]: np.mean(scores_cb) #mean prediction
```

```
Out[24]: 0.8403533513267734
```

```
In [29]: pred_test1
```

```
Out[29]: array([3.16176296, 3.53000451, 3.5517105 , ..., 2.66059136, 4.92398845,
                2.1569985 ])
```

```
In [30]: scores_cb
```

```
Out[30]: [0.8538646099429127,
0.8415649369027963,
0.8243040352648234,
0.8356135322792037,
0.8379186537883418,
0.8396292004435237,
0.8398508071967806,
0.8524512758205053,
0.8346842926926143,
0.8436521689362324]
```

```
In [25]: #submission
final_pred1 = pred_test1/n
ss['default_status'] = final_pred1
ss.to_csv('cb2_blend_submission.csv', index = False)
```

XGBoost MODEL

```
In [26]: scores_xgb = []
pred_test2 = 0
for train_index, test_index in skf.split(X,y):
    print('train index: ', train_index, '\n')
    print('test index: ', test_index, '\n')
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    xgb.fit(X_train, y_train, eval_set= [(X_train, y_train), (X_test, y_test)], verbose=0)
    pred = xgb.predict_proba(X_test)[:,1]
    print('score: ', roc_auc_score(y_test, pred))
    scores_xgb.append(roc_auc_score(y_test, pred))
    test_pred2 = xgb.predict_proba(test)[:,1]
    pred_test2 += test_pred2
```

```
train index: [ 0 1 2 ... 55996 55997 55999]
```

```
test index: [ 11 56 58 ... 55983 55984 55998]
```

```
[0] validation_0-error:0.317837 validation_1-error:0.300714
Multiple eval metrics have been passed: 'validation_1-error' will be used for
early stopping.
```

```
Will train until validation_1-error hasn't improved in 120 rounds.
[100] validation_0-error:0.260099 validation_1-error:0.254286
[200] validation_0-error:0.247619 validation_1-error:0.245179
[299] validation_0-error:0.238651 validation_1-error:0.244107
score: 0.8507847330414031
train index: [ 1 2 4 ... 55997 55998 55999]
```

```
test index: [ 0 3 10 ... 55981 55986 55993]
```

```
[0] validation_0-error:0.315218 validation_1-error:0.315357
Multiple eval metrics have been passed: 'validation_1-error' will be used for
early stopping.
```

```
In [28]: np.mean(scores_xgb)
```

```
Out[28]: 0.8380533211359753
```

```
In [31]: final_pred2 = pred_test2/n
ss['default_status'] = final_pred2
ss.to_csv('xgb2_blend_submission.csv', index = False)
```

LIGHT GRADIENT BOOSTING MODEL

```
In [32]: scores_lgb = []
pred_test3 = 0
for train_index, test_index in skf.split(X,y):
    print('train index: ', train_index, '\n')
    print('test index: ', test_index, '\n')
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    lgb.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_test, y_test)], verbose=0)
    pred = lgb.predict_proba(X_test)[:,1]
    print('score: ', roc_auc_score(y_test, pred))
    scores_lgb.append(roc_auc_score(y_test, pred))
    test_pred3 = lgb.predict_proba(test)[:,1]
    pred_test3 += test_pred3
```

```
train index: [ 0 1 2 ... 55996 55997 55999]
```

```
test index: [ 11 56 58 ... 55983 55984 55998]
```

Training until validation scores don't improve for 120 rounds.

```
[100] training's binary_logloss: 0.365324 valid_1's binary_logloss: 0.394
553
```

```
[200] training's binary_logloss: 0.332051 valid_1's binary_logloss: 0.396
52
```

Early stopping, best iteration is:

```
[106] training's binary_logloss: 0.362855 valid_1's binary_logloss: 0.394
477
```

```
score: 0.8517341188511337
```

```
train index: [ 1 2 4 ... 55997 55998 55999]
```

```
test index: [ 0 3 10 ... 55981 55986 55993]
```

Training until validation scores don't improve for 120 rounds.

```
[100] training's binary_logloss: 0.364357 valid_1's binary_logloss: 0.407
27
```

```
[200] training's binary_logloss: 0.331725 valid_1's binary_logloss: 0.407
499
```

Early stopping, best iteration is:

```
[139] training's binary_logloss: 0.350836 valid_1's binary_logloss: 0.406
752
```

```
score: 0.8386280413504132
```

```
train index: [ 0 1 2 ... 55997 55998 55999]
```

```
test index: [ 4 21 25 ... 55956 55965 55995]
```

Training until validation scores don't improve for 120 rounds.

```
[100] training's binary_logloss: 0.361496 valid_1's binary_logloss: 0.425
859
```

```
[200] training's binary_logloss: 0.328885 valid_1's binary_logloss: 0.426
724
```

Early stopping, best iteration is:

```
[117] training's binary_logloss: 0.355089 valid_1's binary_logloss: 0.425
334
```

```
score: 0.820926122451021
```

```
train index: [ 0 1 2 ... 55996 55997 55998]
```

```
test index: [ 7 17 18 ... 55987 55992 55999]
```

```

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.363579    valid_1's binary_logloss: 0.411
121
[200] training's binary_logloss: 0.330277    valid_1's binary_logloss: 0.411
887
Early stopping, best iteration is:
[110] training's binary_logloss: 0.359771    valid_1's binary_logloss: 0.410
869
score: 0.8356321595237222
train index: [    0    1    2 ... 55997 55998 55999]

test index: [    9   27   28 ... 55969 55977 55978]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.363432    valid_1's binary_logloss: 0.408
97
[200] training's binary_logloss: 0.330487    valid_1's binary_logloss: 0.409
121
Early stopping, best iteration is:
[128] training's binary_logloss: 0.353501    valid_1's binary_logloss: 0.408
32
score: 0.8385650881629207
train index: [    0    1    2 ... 55997 55998 55999]

test index: [    5   46   57 ... 55950 55970 55988]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.363641    valid_1's binary_logloss: 0.408
585
[200] training's binary_logloss: 0.331443    valid_1's binary_logloss: 0.409
906
Early stopping, best iteration is:
[113] training's binary_logloss: 0.358962    valid_1's binary_logloss: 0.408
215
score: 0.8363912077197415
train index: [    0    1    2 ... 55997 55998 55999]

test index: [   14   24   32 ... 55973 55974 55976]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.362926    valid_1's binary_logloss: 0.409
949
Early stopping, best iteration is:
[67] training's binary_logloss: 0.376759    valid_1's binary_logloss: 0.409
462
score: 0.8375198420360171
train index: [    0    1    3 ... 55997 55998 55999]

test index: [    2    6   12 ... 55958 55989 55991]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.364911    valid_1's binary_logloss: 0.397
628
[200] training's binary_logloss: 0.331497    valid_1's binary_logloss: 0.398
31
Early stopping, best iteration is:

```



```

[133] training's binary_logloss: 0.352612    valid_1's binary_logloss: 0.397
454
score: 0.8484394264531059
train index: [    0    2    3 ... 55997 55998 55999]

test index: [    1   13   23 ... 55962 55980 55990]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.362725    valid_1's binary_logloss: 0.412
123
[200] training's binary_logloss: 0.3308      valid_1's binary_logloss: 0.413
291
Early stopping, best iteration is:
[101] training's binary_logloss: 0.362326    valid_1's binary_logloss: 0.412
004
score: 0.8344143306734776
train index: [    0    1    2 ... 55995 55998 55999]

test index: [    8   15   22 ... 55994 55996 55997]

Training until validation scores don't improve for 120 rounds.
[100] training's binary_logloss: 0.363784    valid_1's binary_logloss: 0.406
999
Early stopping, best iteration is:
[75] training's binary_logloss: 0.373766     valid_1's binary_logloss: 0.406
683
score: 0.8396772971250941

```

```
In [33]: np.mean(scores_lgb)
```

```
Out[33]: 0.8381927634346648
```

```
In [34]: final_pred3 = pred_test3/n
ss['default_status'] = final_pred1
ss.to_csv('lgb2_blend_submission.csv', index = False)
```

RANDOMFOREST MODEL

```
In [35]: scores_rfc = []
pred_test4 = 0
for train_index, test_index in skf.split(X,y):
    print('train index: ', train_index, '\n')
    print('test index: ', test_index, '\n')
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    rfc.fit(X_train, y_train)
    pred = rfc.predict_proba(X_test)[: ,1]
    print('score: ', roc_auc_score(y_test, pred))
    scores_rfc.append(roc_auc_score(y_test, pred))
    test_pred4 = rfc.predict_proba(test)[: ,1]
    pred_test4 += test_pred4
```

```
train index: [    0     1     2 ... 55996 55997 55999]
```

```
test index: [   11    56    58 ... 55983 55984 55998]
```

```
score: 0.8465616952448267
```

```
train index: [     1     2     4 ... 55997 55998 55999]
```

```
test index: [     0     3    10 ... 55981 55986 55993]
```

```
score: 0.8358065312293514
```

```
train index: [     0     1     2 ... 55997 55998 55999]
```

```
test index: [     4    21    25 ... 55956 55965 55995]
```

```
score: 0.8198869636711814
```

```
train index: [     0     1     2 ... 55996 55997 55998]
```

```
test index: [     7    17    18 ... 55987 55992 55999]
```

```
score: 0.8293548643583026
```

```
train index: [     0     1     2 ... 55997 55998 55999]
```

```
test index: [     9    27    28 ... 55969 55977 55978]
```

```
score: 0.8345217687810486
```

```
train index: [     0     1     2 ... 55997 55998 55999]
```

```
test index: [     5    46    57 ... 55950 55970 55988]
```

```
score: 0.8346735183463846
```

```
train index: [     0     1     2 ... 55997 55998 55999]
```

```
test index: [    14    24    32 ... 55973 55974 55976]
```

```
score: 0.8327792158896264
```

```
train index: [     0     1     3 ... 55997 55998 55999]
```

```
test index: [     2     6    12 ... 55958 55989 55991]
```

```
score: 0.8435358921917191
```

```
train index: [     0     2     3 ... 55997 55998 55999]
```

```
test index: [     1    13    23 ... 55962 55980 55990]
```

```
score: 0.8292353179276848
train index: [    0     1     2 ... 55995 55998 55999]

test index: [    8    15    22 ... 55994 55996 55997]

score: 0.8341715200068405
```

```
In [36]: np.mean(scores_rfc)
```

```
Out[36]: 0.8340527287646966
```

```
In [37]: final_pred4 = pred_test4/n
ss['default_status'] = final_pred1
ss.to_csv('rfc2_blend_submission.csv', index = False)
```

```
In [38]: #read in the submission files
cb2 = pd.read_csv('/content/cb2_blend_submission.csv')
xgb2 = pd.read_csv('/content/xgb2_blend_submission.csv')
lgb2 = pd.read_csv('/content/lgb2_blend_submission.csv')
rfc2 = pd.read_csv('/content/rfc2_blend_submission.csv')
```

BLEND THE MODELS BASED ON LB PERFORMANCE ¶

```
In [41]: final_blend2 = ((cb2.default_status * 0.7 + xgb2.default_status * 0.3) +
                        (lgb2.default_status * 0.8 + rfc2.default_status * 0.2))
ss['default_status'] = final_blend2
ss.to_csv('final_submission.csv', index = False)
```

```
In [ ]:
```